

OVERTPOLY

AN ALGORITHM FOR FORWARD REACHABILITY ANALYSIS OF NEURAL
FEEDBACK SYSTEMS

SAMUEL (IFEOLUWA) AKINWANDE


STANFORD UNIVERSITY

SAMAKIN@STANFORD.EDU


03/17/2025

INTRODUCTION

Neural networks have found recent success as controllers for dynamical systems



(a) Drone Racing ¹




(b) Autonomous Driving ²

¹Credit: (Kaufmann et al. 2023)


²Credit: (Ettinger et al. 2021)

NEURAL FEEDBACK SYSTEMS

The resulting systems are what we call *Neural Feedback Systems* (NFS)



NEURAL FEEDBACK SYSTEMS



Assume $x \in \mathbb{R}^n$, $f(x) = [f_1(x), \dots, f_n(x)]$, where each $f_i : \mathbb{R}^n \rightarrow \mathbb{R}$, $\epsilon \in E$, π is a neural network

The system (\mathcal{D}) evolves such that for each $i \in [1..n]$


$$next^{\mathcal{D}}(x)_i = \left\{ x_i + (f_i(x) + \pi(o(x))_i + \epsilon) \cdot \delta \mid \epsilon \in E \right\}$$

where δ is the time step size, and E the error set

NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where


- n is the dimensionality of the system



NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where


- n is the dimensionality of the system
- I is the set of initial states



NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where


- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i



NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where


- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i
- E is the bounded error set



NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where


- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i
- E is the bounded error set
- u is the neural network controller



NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where


- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i
- E is the bounded error set
- u is the neural network controller
- δ is the time step size




NEURAL FEEDBACK SYSTEMS

A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where

- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i
- E is the bounded error set
- u is the neural network controller
- δ is the time step size
- T is the time horizon




NEURAL FEEDBACK SYSTEMS



A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where

- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i
- E is the bounded error set
- u is the neural network controller
- δ is the time step size
- T is the time horizon
- G is the set of goal states


NEURAL FEEDBACK SYSTEMS



A NFS is the tuple $\langle n, I, F, E, u, \delta, T, G, A \rangle$, where

- n is the dimensionality of the system
- I is the set of initial states
- F is the set of functions f_i
- E is the bounded error set
- u is the neural network controller
- δ is the time step size
- T is the time horizon
- G is the set of goal states
- A is the set of unsafe states at each time step

REACH-AVOID PROPERTIES




A *trajectory* $\tau^{\mathcal{D}}(\mathcal{X}_0)$ is a sequence of states $(\mathcal{X}_0, \mathcal{X}_1, \dots, \mathcal{X}_T)$, where $\mathcal{X}_0 \subseteq I$, and for each $t \in [1..T]$, $\mathcal{X}_t = \text{next}^{\mathcal{D}}(\mathcal{X}_{t-1})$

A system (\mathcal{D}) is *safe* if for all trajectories $\tau^{\mathcal{D}}(\mathcal{X}_0)$,

$$\forall x_0 \in I. \exists t \in [0..T]. \tau^{\mathcal{D}}(\{x_0\})_t \subseteq G, \quad (1)$$

$$\forall t \in [0..T], \tau^{\mathcal{D}}(I)_t \cap A(t) = \emptyset \quad (2)$$

UNICYCLE CAR MODEL



The state variables are $\mathbf{x} = (x, y, \omega, v)$. We define the neural feedback system \mathcal{U} as follows:

$$\langle 4, I^{\mathcal{U}}, F^{\mathcal{U}}, E^{\mathcal{U}}, u^{\mathcal{U}}, 0.2, 50, G^{\mathcal{U}}, \emptyset \rangle$$

$$I^{\mathcal{U}} = [9.5, 9.55] \times [-4.5, -4.45] \times [2.1, 2.11] \times [1.5, 1.51],$$


$$F^{\mathcal{U}} = (\mathbf{x}_4 \cos(\mathbf{x}_3), \quad \mathbf{x}_4 \sin(\mathbf{x}_3), \quad 0, \quad 0)$$

$$E^{\mathcal{U}} = \{0\} \times \{0\} \times \{0\} \times [-10^{-4}, 10^{-4}],$$


$$G^{\mathcal{U}} = [-0.6, 0.6] \times [-0.2, 0.2] \times [-0.06, 0.06] \times [-0.3, 0.3],$$

$u^{\mathcal{U}}$ is a ReLU neural network with one hidden layer, 500 neurons, and four outputs.


VERIFICATION APPROACHES



(a) Reachability Analysis




(b) Invariance Analysis




(c) Falsification

REACHABILITY APPROACHES




(a) Abstraction Propagation



(b) Combinatorial Optimization

ABSTRACTION PROPAGATION




Tools using these methods propagate an abstract representation of the system to compute reachable sets.

Representations include: Taylor models, Bernstein polynomials, zonotopes, and polytopes.

The CORA tool^a is a representative example of this approach.

^aKochdumper and Althoff 2023.


LIMITATIONS: ABSTRACTION PROPAGATION



Abstraction propagation can be computationally efficient, but the inexactness of the abstraction often leads to excess conservatism

This especially affects systems with nonlinear dynamics, and long time horizons.

COMBINATORIAL OPTIMIZATION




Tools using these methods solve combinatorial problems to compute reachable sets.

They often represent the system as integer programs, hybrid zonotopes, or marching trees.

The OVERTVerify tool^a is a representative example of this approach.

^aSidrane et al. 2022.

LIMITATIONS: COMBINATORIAL OPTIMIZATION



While combinatorial optimization can be arbitrarily precise, computing reachable sets is computationally expensive.

The problem quickly becomes intractable when the system dynamics are nonlinear, or the time horizon is long.

OVERTPOLY

OvertPoly is a **combinatorial algorithm** for forward reachability analysis of Neural Feedback Systems with **computational efficiency** comparable to abstraction propagation methods.

OVERTPOLY

Our contributions are:

- We introduce a **novel combinatorial abstraction** for nonlinear dynamical systems

OVERTPOLY

Our contributions are:

- We introduce a **novel combinatorial abstraction** for nonlinear dynamical systems
- Using our abstraction, we define an **efficient representation** of nonlinear neural feedback systems

OVERTPOLY

Our contributions are:

- We introduce a **novel combinatorial abstraction** for nonlinear dynamical systems
- Using our abstraction, we define an **efficient representation** of nonlinear neural feedback systems
- We use this representation to define novel algorithms for forward reachability analysis

OVERTPOLY

Our contributions are:

- We introduce a **novel combinatorial abstraction** for nonlinear dynamical systems
- Using our abstraction, we define an **efficient representation** of nonlinear neural feedback systems
- We use this representation to define novel algorithms for forward reachability analysis
- We demonstrate an **order of magnitude improvement** in performance compared to the current state-of-the-art

POLYHEDRA

Definition (k -Simplex)

A polyhedron \mathcal{S} is a k -simplex if it is the convex hull of $k + 1$ affinely independent points in \mathbb{R}^k . A polyhedron is a simplex if it is a k -simplex for some k , and k is called its *dimension*.

POLYHEDRA

Definition (Simplicial k -Complex)

A *simplicial complex* \mathcal{C} is a set of simplices such that:

- Every face of a simplex in \mathcal{C} is also in \mathcal{C}
- Every non-empty intersection of two simplices $\mathcal{S}_1, \mathcal{S}_2 \in \mathcal{C}$ is a face of both \mathcal{S}_1 and \mathcal{S}_2

POLYHEDRA

Definition (Point Set Triangulation)

If P is a finite set of points in \mathbb{R}^n , then a pure simplicial n -complex \mathcal{C} is a *point set triangulation of P* if $P = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathbf{vert}(\mathcal{S})$ and $\mathbf{conv}(P) = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathcal{S}$.

POLYHEDRA

Definition (Point Set Triangulation)

If P is a finite set of points in \mathbb{R}^n , then a pure simplicial n -complex \mathcal{C} is a *point set triangulation of P* if $P = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathbf{vert}(\mathcal{S})$ and $\mathbf{conv}(P) = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathcal{S}$.

Let C be a closed n -ball. We call C^O its open n -ball, and C^S the hypersphere forming its surface. We call $V_P(C) = C \cap P$ the vertices of C .

POLYHEDRA

Definition (Point Set Triangulation)

If P is a finite set of points in \mathbb{R}^n , then a pure simplicial n -complex \mathcal{C} is a *point set triangulation of P* if $P = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathbf{vert}(\mathcal{S})$ and $\mathbf{conv}(P) = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathcal{S}$.

Let C be a closed n -ball. We call C^O its open n -ball, and C^S the hypersphere forming its surface. We call $V_P(C) = C \cap P$ the vertices of C .

If \mathcal{C} is a point set triangulation of P , and \mathcal{S} is a n -simplex in \mathcal{C} , then $C(\mathcal{S})$ is the smallest closed n -ball C such that $\mathcal{S} \subseteq C$ and $C^O \cap \mathbf{vert}(\mathcal{S}) = \emptyset$.

POLYHEDRA

Definition (Point Set Triangulation)

If P is a finite set of points in \mathbb{R}^n , then a pure simplicial n -complex \mathcal{C} is a *point set triangulation of P* if $P = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathbf{vert}(\mathcal{S})$ and $\mathbf{conv}(P) = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathcal{S}$.

Let C be a closed n -ball. We call C^O its open n -ball, and C^S the hypersphere forming its surface. We call $V_P(C) = C \cap P$ the vertices of C .

If \mathcal{C} is a point set triangulation of P , and \mathcal{S} is a n -simplex in \mathcal{C} , then $C(\mathcal{S})$ is the smallest closed n -ball C such that $\mathcal{S} \subseteq C$ and $C^O \cap \mathbf{vert}(\mathcal{S}) = \emptyset$.

\mathcal{S} satisfies the *Delaunay condition* and is called a *Delaunay simplex of P* if $V_P(C(\mathcal{S})^O) = \emptyset$.

POLYHEDRA

Definition (Point Set Triangulation)

If P is a finite set of points in \mathbb{R}^n , then a pure simplicial n -complex \mathcal{C} is a *point set triangulation of P* if $P = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathbf{vert}(\mathcal{S})$ and $\mathbf{conv}(P) = \bigcup_{\mathcal{S} \in \mathcal{C}} \mathcal{S}$.

Let C be a closed n -ball. We call C^O its open n -ball, and C^S the hypersphere forming its surface. We call $V_P(C) = C \cap P$ the vertices of C .

If \mathcal{C} is a point set triangulation of P , and \mathcal{S} is a n -simplex in \mathcal{C} , then $C(\mathcal{S})$ is the smallest closed n -ball C such that $\mathcal{S} \subseteq C$ and $C^O \cap \mathbf{vert}(\mathcal{S}) = \emptyset$.

\mathcal{S} satisfies the *Delaunay condition* and is called a *Delaunay simplex of P* if $V_P(C(\mathcal{S})^O) = \emptyset$.

\mathcal{C} is a *Delaunay triangulation* if every n -simplex in \mathcal{C} satisfies the Delaunay condition.

POLYHEDRA

Definition (Bounding Set)

A *bounding set* is a tuple $\mathcal{B} = \langle n, P, L, U \rangle$, where $n \in \mathbb{N}$, P is finite a set of points in \mathbb{R}^n , and L and U are functions from P to \mathbb{R} , such that $L(p) \leq U(p)$ for all $p \in P$. The *domain* of \mathcal{B} is defined as $\mathbf{dom}(\mathcal{B}) = \mathbf{conv}(P)$.

POLYHEDRA

Definition (Polyhedron formed by Bounding Set)

Let $\mathcal{B} = \langle n, P, L, U \rangle$ be a bounding set. We define the *vertices* of the bounding set as:

$$V(\mathcal{B}) := \{(p, L(p)) : p \in P\} \cup \{(p, U(p)) : p \in P\}.$$

We define the $(n + 1$ -dimensional) *polyhedron formed by \mathcal{B}* as

$$\mathcal{P}(\mathcal{B}) := \mathbf{conv}(V(\mathcal{B})).$$

POLYHEDRAL ENCLOSURES

Definition (Polyhedral Enclosure)

Let $\mathcal{B} = \langle n, P, L, U \rangle$ be a bounding set, let Δ be a Delaunay triangulation of P , and let Δ_n be the set of all n -simplices in Δ . We then define the bounding set associated with a simplex $\mathcal{S} \in \Delta_n$ as:

$$\mathcal{B}_{\mathcal{S}} := \langle n, \mathbf{vert}(\mathcal{S}), L^{\mathbf{vert}(\mathcal{S})}, U^{\mathbf{vert}(\mathcal{S})} \rangle,$$

We define the *polyhedral enclosure* formed by \mathcal{B} and Δ as:

$$\mathcal{E}(\mathcal{B}, \Delta) := \bigcup_{\mathcal{S} \in \Delta_n} \mathcal{P}(\mathcal{B}_{\mathcal{S}}).$$

POLYHEDRAL ENCLOSURES

Definition (Bounding Set Composition)

Let $\mathcal{B}^f = \langle n, P, L^f, U^f \rangle$ and $\mathcal{B}^g = \langle n, P, L^g, U^g \rangle$ be bounding sets with the same dimension and over the same set of points P . For $\bowtie \in \{+, -, \times, /\}$, we define $\mathcal{B}^f \bowtie \mathcal{B}^g := \langle n, P, L, U \rangle$, where, for all $\mathbf{x} \in P$,

$$L(\mathbf{x}), U(\mathbf{x}) = \begin{cases} L^f(\mathbf{x}) + L^g(\mathbf{x}), U^f(\mathbf{x}) + U^g(\mathbf{x}) & \text{if } \bowtie = +, \\ L^f(\mathbf{x}) - L^g(\mathbf{x}), U^f(\mathbf{x}) - U^g(\mathbf{x}) & \text{if } \bowtie = -, \\ \min(Bounds), \max(Bounds) & \text{if } \bowtie \in \{\times, \div\}, \end{cases}$$

where $Bounds = \{h_1(x) \bowtie h_2(x) \mid h_1 \in \{L^f, U^f\}, h_2 \in \{L^g, U^g\}\}$ is the set of potential bounds when using multiplication or division.

REFERENCES

- Ettinger, Scott et al. (2021). “Large scale interactive motion forecasting for autonomous driving: The waymo open motion dataset”. In: *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 9710–9719.
- Kaufmann, Elia et al. (2023). “Champion-level drone racing using deep reinforcement learning”. In: *Nature* 620.7976, pp. 982–987.
- Kochdumper, Niklas and Matthias Althoff (2023). “Constrained polynomial zonotopes”. In: *Acta Informatica* 60.3, pp. 279–316.
- Sidrane, Chelsea et al. (2022). “Overt: An algorithm for safety verification of neural network control policies for nonlinear systems”. In: *Journal of Machine Learning Research* 23.117, pp. 1–45.